

A FREE WORKBOOK FROM IMPACT LOOP

Nonprofit AI Adoption *Workbook.*

*A 30-day, self-directed guide to running
one ADAPT cycle on your team's real work,
without losing the mission.*

BUILT BY

Rovonn Russell
Impact Loop, Toronto
May 2026

IMPACTLOOP.CA

A working tool, *not a strategy doc.*

By the end of 30 days, your team will have audited your real communications work, identified one specific place AI genuinely fits, built a small working system, run it through the Trust Test, and handed it off as a runbook the team can maintain.

The pace is one phase per week. Week 1 is Assess. Week 2 is Discover. Week 3 is Apply. Week 4 is Produce and Transform. Each phase has a worksheet at the end. Do the worksheet before moving on.

You can run this self-directed, or with facilitation. **ADAPT AI Training** is the 4-week paid version we deliver at Impact Loop, where we run the same cycle with your team and ship a system you can run without us. Pricing typically lands \$5,000 to \$15,000 depending on team size. Discovery calls are free.

For now, what you need: one team member willing to be the operator on this cycle, two hours a week of their time, and the willingness to slow down before you reach for a tool.

"The orgs that compound on AI adoption are the ones that keep running ADAPT. Not the ones that buy the most tools."

Assess: *what is this work actually doing?*

The hardest part of nonprofit AI adoption is not technical. It is honest.

Most nonprofits skip Assess because the question feels obvious. Comms is comms, right? Reports are reports. Donor emails are donor emails. We know what we do.

You don't, actually. Not in the resolution that matters for AI decisions. The Assess phase forces a higher resolution: every recurring piece of communication your org produces gets sorted into one of three jobs. The job determines whether AI belongs anywhere near it.

The three jobs

Job 1, Operational

The reader doesn't care who wrote it. They care that it is accurate and clear. Internal team docs. FAQ pages. Meeting notes. Vendor briefs. Process documentation. Onboarding emails for new contractors. Foundation acknowledgement receipts. SOPs. AI is good at this work. Sometimes better than a tired human who would otherwise rush it.

Job 2, Informational

The reader cares the information is correct and the explanation is clear. They mostly do not care who wrote it. Blog posts explaining a program. How-to guides for community members. Plain-language summaries of research. Press release drafts. Service descriptions on the website. AI can do most of the work here, but a human should still rewrite the voice-heavy parts. Use AI as co-writer, not as the writer.

Job 3, Relational

The reader is paying attention to whether the voice still sounds like a person who cares. Donor appeal letters. Grant reports to your largest funders. The story of why your org exists. Sensitive community communications. Personal thank-you notes. The newsletter that goes

to your community. **AI cannot do this work.** Not because the model isn't smart enough, because the reader can tell. The second they can tell, the trust is gone.

Why the sort matters

The nonprofits that get AI adoption right run AI on Operational and most of Informational. They do not let AI touch Relational. Ever. Even when the team feels overwhelmed and someone says "*let's just use ChatGPT for the year-end appeal.*" That is the moment AI breaks the brand.

Most nonprofits running AI badly right now have it backwards. They feel most overwhelmed by the Relational work, so they reach for AI exactly where it cannot help. The team feels productive for two weeks, donations soften, and nobody connects the dots until the next quarterly review.

Sort the work first. Decide where AI is allowed second.

WORKSHEET — WEEK 1

Sort your work

List every recurring piece of communication your org produces. Go through the last 90 days of email outputs, blog posts, social, reports, donor mail, board memos. Aim for at least 25 items.

COMMUNICATION	FREQUENCY	AUDIENCE	JOB
<i>Year-end donor appeal email</i>	<i>Annual</i>	<i>Donors</i>	<i>Relational</i>
<i>M&E summary in funder report</i>	<i>Quarterly</i>	<i>Foundation officers</i>	<i>Relational</i>
<i>Service description on website</i>	<i>Ongoing</i>	<i>Search visitors</i>	<i>Informational</i>
<i>Internal weekly team update</i>	<i>Weekly</i>	<i>Internal</i>	<i>Operational</i>
_____	_____	_____	_____
_____	_____	_____	_____
_____	_____	_____	_____
_____	_____	_____	_____
_____	_____	_____	_____
_____	_____	_____	_____

When you finish: count how many rows are in each column. Most nonprofits are surprised. The honest sort is the whole game.

Discover: *where does AI actually fit?*

You have a sorted inventory. Week 2 narrows it. Pick one specific outcome you will build toward.

Five places AI consistently fits

1. Inbox triage

A small AI agent that classifies incoming email and routes it. No drafting, no replying, just routing. Saves the executive director or operations lead 4-7 hours a week. Low risk because nothing goes out.

2. Document summarization

Foundation reports, board packets, research. AI summarization is genuinely good at compressing structured documents into the 3-5 bullet points the team actually needs. Use this for incoming documents, not outgoing.

3. Internal documentation drafting

SOPs, onboarding docs, runbooks, training materials. AI produces structured first drafts in minutes. The team rewrites the parts that need org-specific context.

4. Data extraction from PDFs and CRM exports

Pulling stats, dates, names, amounts out of structured documents into spreadsheet form. Classification work. Deterministic-shaped. Low risk if validated.

5. First-pass grant research

AI scans grant databases, summarizes fit, produces a shortlist. The decision-maker still evaluates fit and writes the proposal. This is research-acceleration, not application-writing.

Five places AI does *not* fit

1. Donor relationship layer

Personal thank-you notes. Appeal emails. The story of a beneficiary. The voice has to be yours.

2. Sensitive community communications

Conversations with families in crisis. Communications to youth in the program. Anything in the community-organizer lane where tone-deafness has real consequences.

3. Brand-defining narrative

The "about us" page, the founder's story, the mission statement, the homepage hero copy. These are the places people decide whether to trust you. Write them yourself.

4. Sensitive grant work

Multi-year capacity grant reports, foundation transition narratives, anything where the funder is reading for nuance and partnership signal. Generic phrasing kills next-cycle funding faster than missing a deadline.

5. Stakeholder communication during a crisis

Reputational issue, leadership transition, program closure. AI generates the wrong texture for these moments. Slow down, write it yourself, send it.

WORKSHEET — WEEK 2

Pick your one outcome

From the inventory in Week 1, pick the single workflow you will run AI on in Week 3.

Workflow name:

Job category: ☐ Operational ☐ Informational

Current time per cycle: _____ **How often it runs:**

Owner of the workflow today: _____

Owner of the AI version going forward: _____

Time-savings target: _____

Trust risk if it goes wrong once: ☐ Low ☐ Medium ☐ High

Human review gate before output is used:

Why this workflow and not another:

*If you cannot fill in every row above, you do not have a workflow ready for Week 3.
Pick a different one.*

Apply: *build the smallest version that could work.*

Most nonprofit AI adoption fails here, not because the tools don't work, but because teams overscope the first version.

Two failure modes to avoid

The committee

Five people, three meetings, no shipped version. Apply is owned by **one person**. That person decides what to build, builds it, and reports back. Other team members get involved only at the review gate.

The vendor

A SaaS tool gets bought. Onboarding takes three weeks. Customization eats two more. Output is generic. Quietly stops getting used. Apply uses the simplest tool that could possibly work, usually a generic chat interface (Claude, ChatGPT, or whichever your org has access to) with a saved prompt template, not a category-specific SaaS.

Three sample build patterns

Pattern A · Saved prompt template

For document summarization, drafting, or categorization. The operator writes a prompt that takes a structured input and produces a structured output. Lives in a shared doc. Run manually at first, refine over the week. **Build time: 1-3 hours. Tool cost: ~\$20/month.**

Pattern B · Routing automation

For inbox triage. The operator wires Gmail / Outlook to a no-code platform (Zapier, Make.com, n8n) that runs a classification prompt and routes to the right channel. **Build time: 4-8 hours. Tool cost: ~\$30-100/month including AI.**

Pattern C · Spreadsheet-bound workflow

For data extraction or research. The operator builds a Google Sheet with a column that triggers an AI prompt on each new row. Inputs go in one column, outputs in the next. **Build time: 2-5 hours. Tool cost: ~\$20/month.**

What to log

The v1 must produce its work AND a log of what it did. Date. Input. Output. Time elapsed. Whether the output was used as-is, edited, or thrown away. **Without this log, you have no way to compare with-AI to without-AI honestly.**

WORKSHEET — WEEK 3

Scope the v1

Build pattern: ☐ A ☐ B ☐ C ☐ Other: _____

Tool:

Operator:

Build deadline: _____ **First production run date:**

Where the log lives:

Sample input to test against:

Expected output for that sample:

What "working" looks like at end of Week 3:

If the v1 is not running by end of Week 3, do not advance. Either rescope down or postpone the cycle.

Produce + Transform: *compare honestly, hand off cleanly.*

Run the comparison that determines whether the v1 is worth keeping. Convert what works into a system the team owns.

Produce: the three measurements

Run the v1 alongside the manual process for 2-4 weeks. Track three things:

Time

How long does the workflow take with AI, end-to-end including review? How long did it take without AI? **Be honest about review time.** AI that saves 30 minutes of drafting but adds 45 minutes of review is a net loss.

Quality

Is the output as good as what the team was producing manually? This is where you run **the Trust Test**. A v1 that fails the Trust Test for any output category that touches external audiences gets retired.

Team energy

Does the operator feel more capable or more drained? AI that technically saves time but burns out the operator is a net loss. Ask the operator at the end of each week. Their answer is the data.

A v1 that wins on all three graduates to Transform. A v1 that fails on any of the three either gets redesigned or retired.

The Trust Test

Take a sample of the AI's output. Hand it to one person from your audience without telling them how it was produced. Ask them whether anything reads off.

For Operational workflows

"Audience" means a different team member or a partner. **The bar is lower:** does the output communicate clearly?

For Informational workflows

"Audience" means a community member or stakeholder who reads your content. **The bar is medium:** does the output sound like the org's voice?

For workflows that touch Relational work (rare)

"Audience" means a donor, board member, or major stakeholder. **The bar is high:** do they spot the AI?

If they spot the AI, the output is not ready for that workflow. No exceptions.

Anti-patterns to retire

- **Single-person knowledge.** If only one person knows how to run it, document and share or retire.
- **No review gate.** AI-without-review breaks trust the fastest. Add even a 30-second eyeball check before external send.
- **Same output regardless of input.** Sign that the prompt is too generic. Real workflows produce different outputs for different inputs.
- **Silent abandonment.** Ask the operator directly: "Are you still using this?" If "not really," retire on the spot.

Transform: hand off cleanly

A v1 that survives Produce becomes a system the team owns. Document four things:

1. The runbook

Step-by-step documentation. If a new team member couldn't run the workflow from this document, the runbook isn't done.

2. The maintenance plan

What happens when the AI tool changes pricing, the prompt stops working, the API key rotates, or the operator leaves. One page.

3. The tone-of-voice gate

Specifically for any Informational workflow that touches external audiences. A documented checklist of what the output must sound like before it gets sent.

4. The budget

What does this workflow cost per month in tool fees, API calls, and team time? Knowing the number protects you when the next budget cycle starts.

WORKSHEET — WEEK 4

The hand-off

Runbook lives at:

Maintenance plan owner:

Tone-of-voice gate reviewed by:

Monthly cost (tools + AI + team time): \$ _____

Date the system was officially handed off: _____

Person responsible going forward:

Person who reviews the output before external send:

When all four are filled in and the operator confirms they can run the system without you, the cycle is complete.

WHAT TO DO NEXT

One cycle done. *Pick the next one.*

You have run one ADAPT cycle. The next cycle picks one more workflow from your Week 1 inventory and runs the same five phases.

Most nonprofits we work with run their first three cycles in roughly the same order:

Cycle 1

An Operational workflow (often inbox triage or document summarization). Builds team confidence. Low risk. Genuine time savings.

Cycle 2

An Informational workflow (often internal docs drafting or research-acceleration). Demonstrates that AI can support content work without compromising voice.

Cycle 3

A higher-leverage Operational workflow (often donor data hygiene, M&E reporting, or grant tracking).

After three cycles, most teams stop needing facilitation. They have the framework, they have a few working systems, and they pick the next workflow themselves.

Run the next cycle with us

If you want to run the cycles with experienced facilitation rather than self-directed, ADAPT AI Training is our 4-week paid engagement. We run one full cycle with your team, ship a system you own, and hand off a runbook. Engagements typically range from \$5,000 to \$15,000 depending on team size and scope.

Free discovery call: impactloop.ca/booking

Quick-reference *templates.*

Sample Operational prompt — meeting notes summarizer

You are summarizing internal meeting notes for [ORG NAME].
The audience is the leadership team.
The output goes in the weekly internal update.

Input: meeting transcript pasted below
Output: structured summary with three sections:
1. Decisions made (bullet list)
2. Action items with owner and deadline (table)
3. Open questions for next meeting (bullet list)

Keep the tone direct. Skip filler. If a decision is unclear, flag it as "needs follow-up" rather than guessing.

Transcript: [PASTE]

Sample Informational prompt — program FAQ drafter

You are drafting an FAQ section for [ORG NAME]'s [PROGRAM NAME] page.
The audience is community members considering applying to the program.

Input: program description and existing comms about the program (below)
Output: 8-10 plain-language FAQ items with direct answers.

Voice: warm, clear, no jargon. Avoid corporate-speak. Avoid superlatives.
Match the existing comms voice, read what's pasted before writing.

If you don't have enough information to answer accurately, write
[NEEDS CONFIRMATION] in place of the answer.

Inputs: [PASTE]

Tone-of-voice gate — sample checklist

Before any AI-assisted output goes external, run through:

- ☐ Does this sound like our org, or like every other org?
- ☐ Are there any specific phrases the team would never use?
- ☐ Does it name specific things (people, places, programs, dates) where it should?
- ☐ Does it avoid superlatives that aren't earned?
- ☐ Would a long-time supporter recognize this as ours?

If any answer is no, edit before sending.

One-page maintenance plan template

Workflow:

Tool:

Operator: _____ **Backup operator:** _____

Monthly cost: \$ _____ **Renewal cycle:** _____

Quarterly review date:

What to do if the prompt stops working:

What to do if the operator leaves:

Free workbook. No upsell. The next cycle is the only thing that matters.

Built by Rovonn Russell at Impact Loop · Toronto · May 2026

impactloop.ca · @rovonnrussell